

Q-value methods for reinforcement learning

- 1) SARSA
- 2) Q-learning

Value functions


- To learn an optimal policy (learn how to act), we need *value functions*
- Two types of value functions in RL:
 - **state value function**, denoted $V(s)$
 - **state-action value function**, denoted $Q(s,a)$ ← Our focus today

Bellman equation:

$$Q^{\pi}(s, a) = \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma \sum_{a'} \pi(s', a') Q^{\pi}(s', a') \right]$$

Q-value methods

Learning algorithm:

- Initialize $Q(s, a) = 0$ for all (s, a) pairs
- Repeat **episodes** (e.g. “games”)
 - select s, a
 - Repeat **steps within episode**
 - * execute action a , observe r, s'
 - * select a' based on $Q(s', a^*)$ ϵ -greedy over a^*
 - * update $Q(s, a)$ 
 - * $s = s', a = a'$
 - Until s terminal (where $Q(s', a') = 0$)

Update rule
differs for SARSA
and Q-learning

Q-value update rules

SARSA update rule:

$$\Delta Q_t(s_t, a_t) = \alpha[r_{t+1} + \gamma Q_t(s_{t+1}, a_{t+1}) - Q_t(s_t, a_t)] \quad \textit{On-policy}$$

Q-learning update rule:

$$\Delta Q_t(s_t, a_t) = \alpha[r_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t)] \quad \textit{Off-policy}$$

Tic-tac-toe (in Python)

- $9^3 = 729$ states
 - 9 board positions
 - Each position has 1 of 3 values: {-, 0, X}
 - States are represented by a string: "0----X----"
- 9 actions (at most)
 - Actions are represented by a tuple: (row, col)
- Q values
 - Q is a dictionary of dictionaries
 - $Q[a][s] \leftarrow$ index value for action **a**, state **s**
 - Initialize all values to 0

	0	1	2
0	0	-	-
1	-	X	-
2	-	-	-